



Chap.8 Memory Management

Young Ik Eom
DCLab@SKKU



Backgrounds

□ Types of memories in computer systems

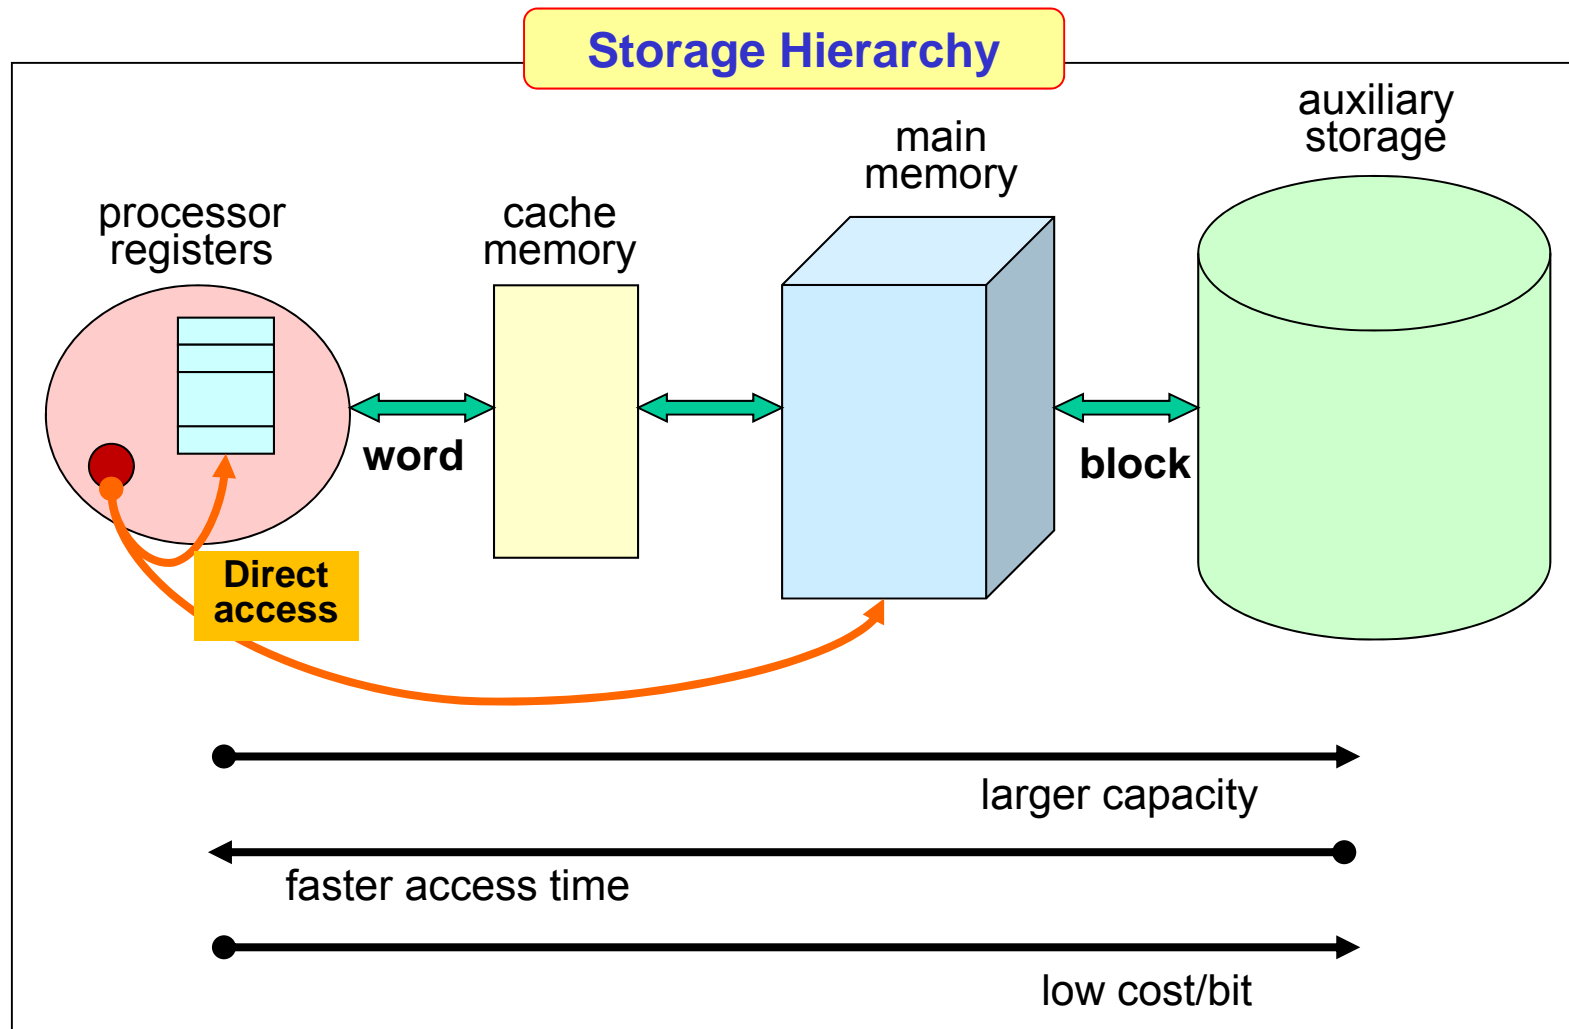
- Processor registers
- Cache memory
- **Main memory**
- **Auxiliary memory**

□ Notes

- Block
 - ✓ Data transfer unit between primary memory and secondary storage
 - ✓ Size: 1 ~ 4 KB (128B ~ 8MB)
- Word
 - ✓ Data transfer unit between primary memory and CPU
 - ✓ Size: 16 ~ 64 bits



Backgrounds





Backgrounds

❑ Cache

❑ CPU register access

- ✓ Generally takes 1~30 cycles of the CPU clock

❑ Memory access

- ✓ Generally takes 50~200 cycles of the CPU clock
- ✓ CPU normally needs to stall during the memory access
- ✓ Intolerable because of the frequency of memory accesses

❑ Cache memory

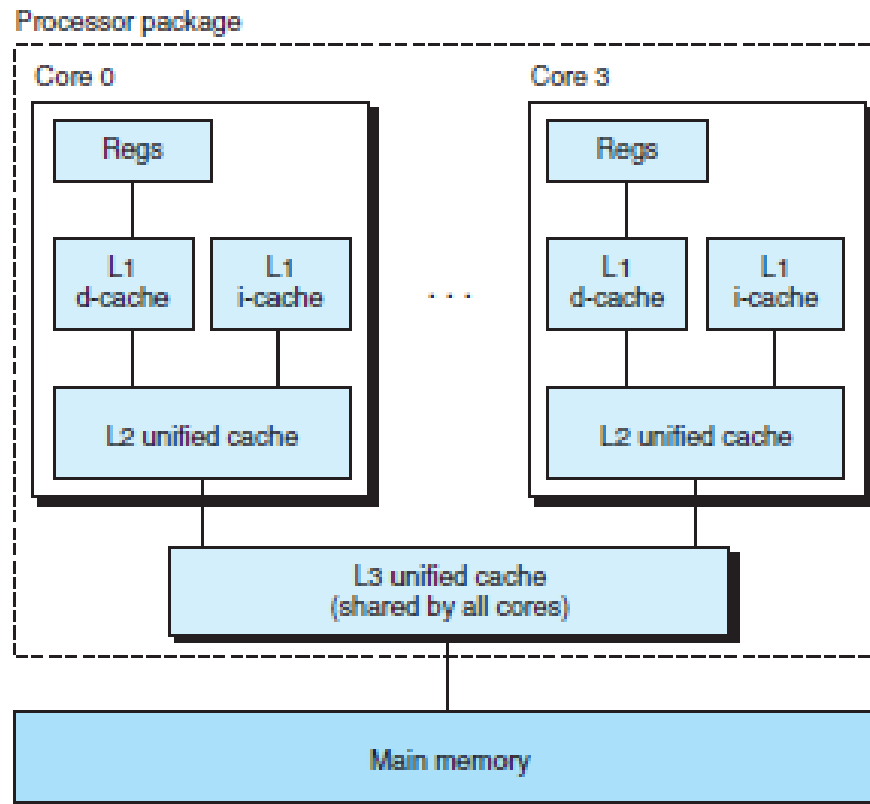
- ✓ Used to accommodate the speed differential



Backgrounds

□ Cache

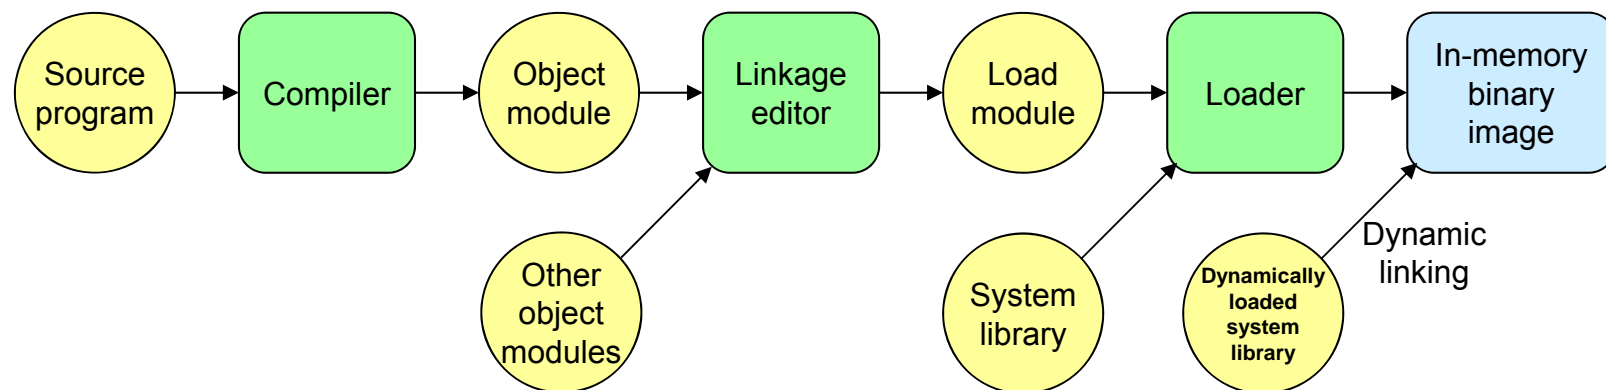
□ Intel Core i7 cache





Backgrounds

□ Address binding





Backgrounds

□ Address binding

□ Compile time binding

- ✓ When it is known at compile time where the process will reside in memory, then **absolute code** can be generated
- ✓ Changing the starting location requires recompilation
- ✓ MS-DOS .COM-format programs



Backgrounds

□ Address binding

□ Load time binding

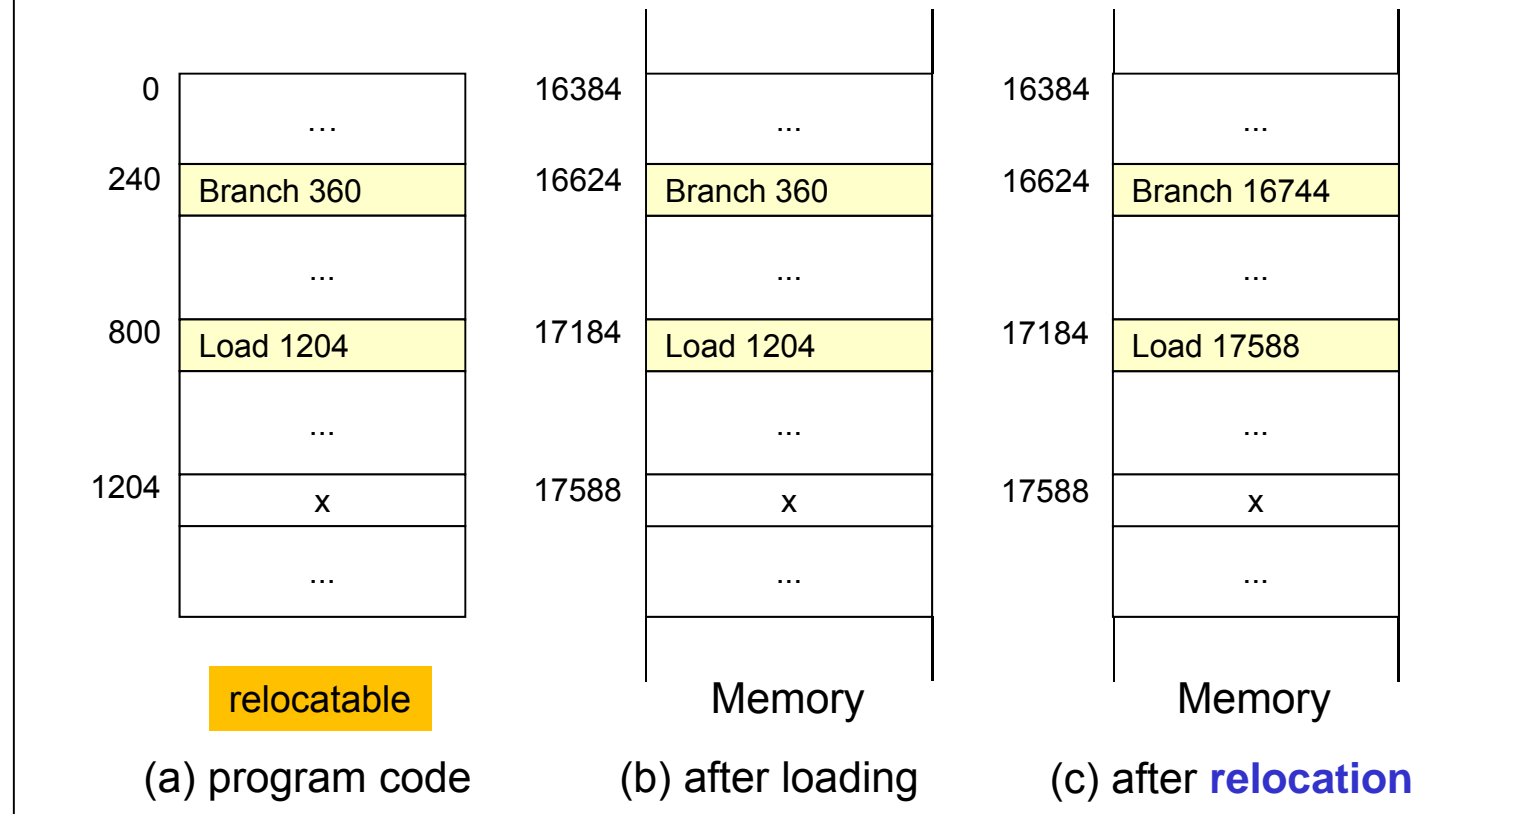
- ✓ When it is not known at compile time where the process will reside in memory, then the compiler must generate **relocatable code**
- ✓ Final binding is delayed until load time
- ✓ Changing the starting location requires reloading or **relocation** of the user code



Backgrounds

Load-time binding

(Allocation address = 16384)





Backgrounds

□ Address binding

□ Run time binding

- ✓ Processes can be moved during its execution from one memory segment to another
- ✓ Special hardware is required
- ✓ Used in most general-purpose operating systems

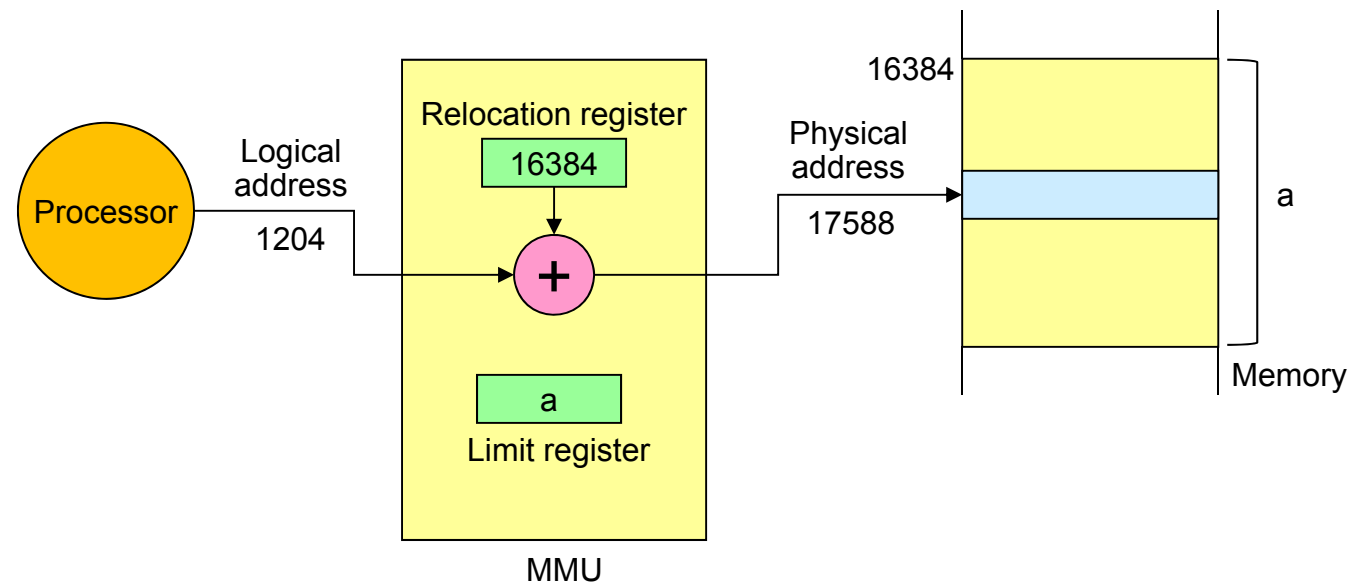


Backgrounds

□ Address binding

□ Run time binding

- ✓ Run-time mapping from logical(virtual) address to physical address is performed by a hardware device called MMU(Memory Management Unit)





Backgrounds

□ Address binding

□ Notes

- ✓ Logical address (virtual address)
 - An address generated by the CPU
- ✓ Physical address
 - An address seen by the memory unit
 - An address loaded into MAR
- ✓ Relocation register (base register)
 - Holds the allocation address (smallest physical address)
- ✓ Limit register
 - Contains the range of logical addresses



Backgrounds

❑ Dynamic loading

- ❑ All routines are kept on disk in a relocatable load format
- ❑ A routine is not loaded until it is called
 - ✓ When a routine needs to call another routine, the caller first checks to see whether the callee has been loaded
 - ✓ Calls the **relocatable linking loader** to load the callee into memory, if necessary
- ❑ Advantages
 - ✓ Better memory space utilization
 - ✓ Unused routine is never loaded
 - ✓ Does not require special support from OS



Backgrounds

❑ Dynamic linking

- ❑ Linking is postponed until execution time
- ❑ Usually used with system libraries
 - ✓ Without this facility, each executable on a system must include a copy of its library
 - Wastes both disk space and main memory



Backgrounds

❑ Dynamic linking

❑ Uses **stub** concept

- ✓ Included in the image for each library routine reference
- ✓ Small piece of code that indicates how to locate or load the appropriate library routine (in memory or from the library)
- ✓ Replaces itself with the address of the routine and executes the routine
 - At the next time that the code segment is reached, the library routine is executed directly without further dynamic linking



Backgrounds

❑ Dynamic linking

❑ Code sharing

- ✓ All processes that use a library execute only one copy of the library code

❑ Supports the concept of shared libraries


❑ Requires help from the OS



Backgrounds

□ Overlay structure

- Keep in memory only those instructions and data that are needed at any given time
- Example) Assembler
 - ✓ Symbol table(20KB), Common routines(30KB), Pass-1(70KB), Pass-2(80KB)
 - ✓ Primary memory size: 150KB

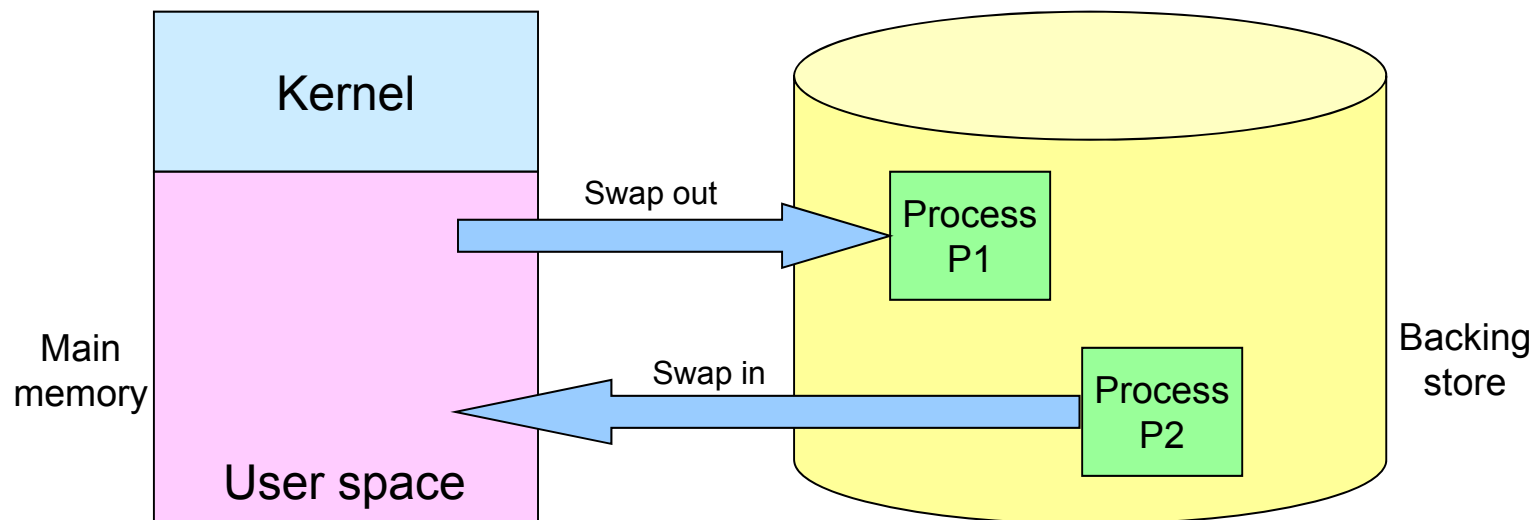




Swapping

❑ Swapping

- ❑ A process can be swapped temporarily out of memory to a **backing store (swap device)**





Swapping

□ Notes on swapping

□ Time quantum vs swap time

- ✓ Time quantum should be substantially larger than swap time (context switch time) for efficient CPU utilization

□ Memory areas to be swapped out

- ✓ Swap only what is actually used

□ Pending I/O

- ✓ If the I/O is asynchronously accessing the user memory for I/O buffers, then the process cannot be swapped

✓ Solutions

- Never swap a process with pending I/O
- Execute I/O operations only into kernel buffers (and deliver it to the process memory when the process is swapped in)



Swapping

□ Notes on swapping

□ Swap device (swap file system)

- ✓ Swap space is allocated as a chunk of disk
 - Contiguous allocation of the process image
 - Fast access time

■ Ordinary file system

- Discontiguous allocation of the file data blocks
- Focuses on the efficient file system space management



Contiguous Memory Allocation



Contiguous Memory Allocation

❑ Basic policies

- ❑ Each process (context) is contained in a single contiguous section of memory

❑ Policies for memory organization

- ❑ Number of processes in memory
 - ✓ Affects multiprogramming degree
- ❑ Amount of memory space allocated for each process
- ❑ Memory partition methods
 - ✓ Fixed(static) partition multiprogramming
 - ✓ Variable(dynamic) partition multiprogramming



Contiguous Memory Allocation

- ❑ Uniprogramming
- ❑ Fixed partition multiprogramming
- ❑ Variable partition multiprogramming

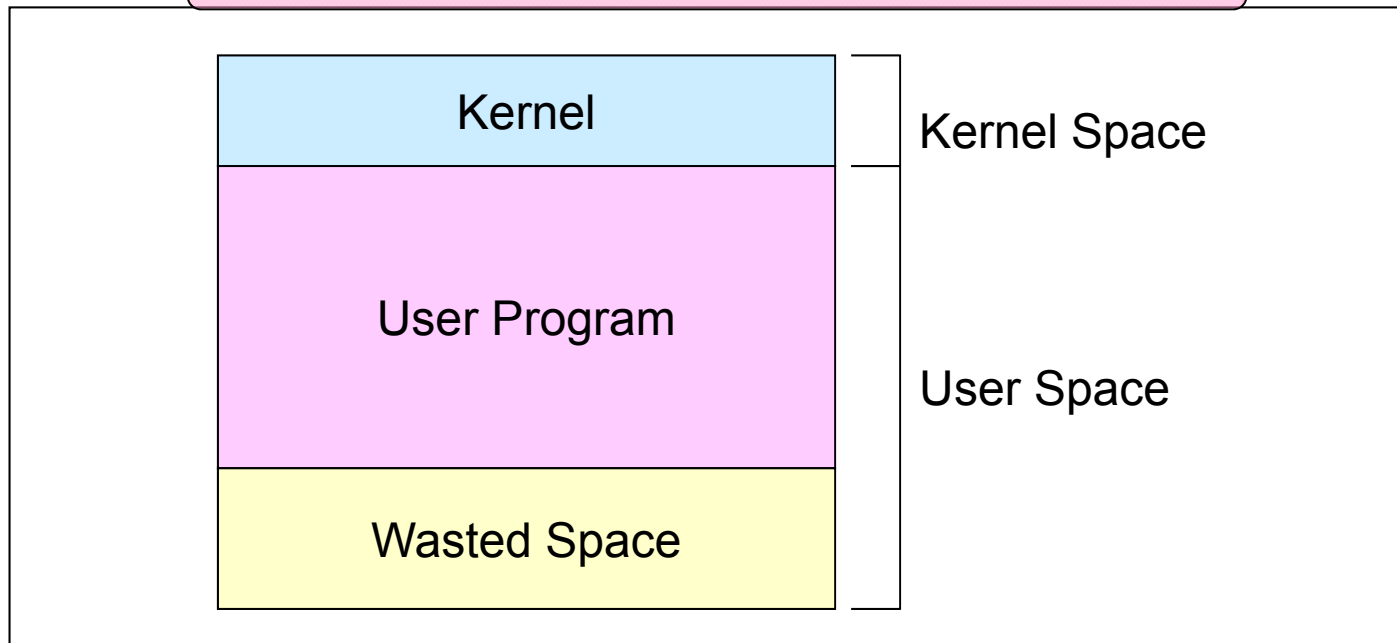


Uniprogramming

□ Uniprogramming system

- Only 1 process in memory
- Simple memory management scheme

Memory state in uniprogramming systems





Uniprogramming

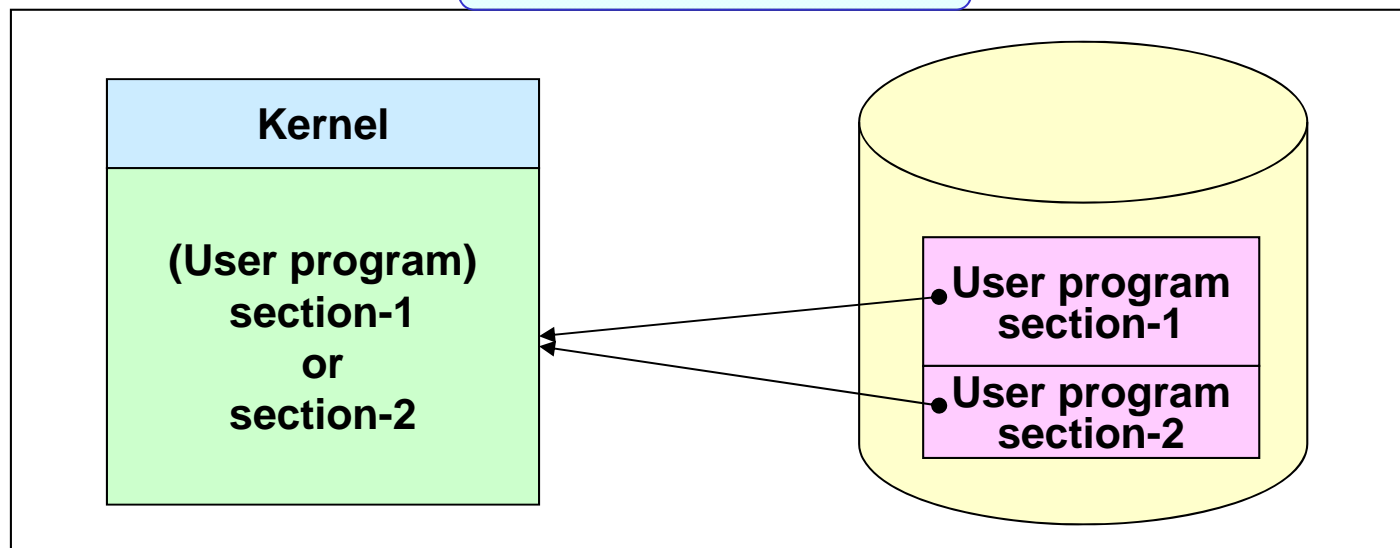
❑ Issue-1

❑ Program-size > memory-size

✓ Uses **overlay structure**

✓ Requires support from compiler/linker/loader

Overlay structure



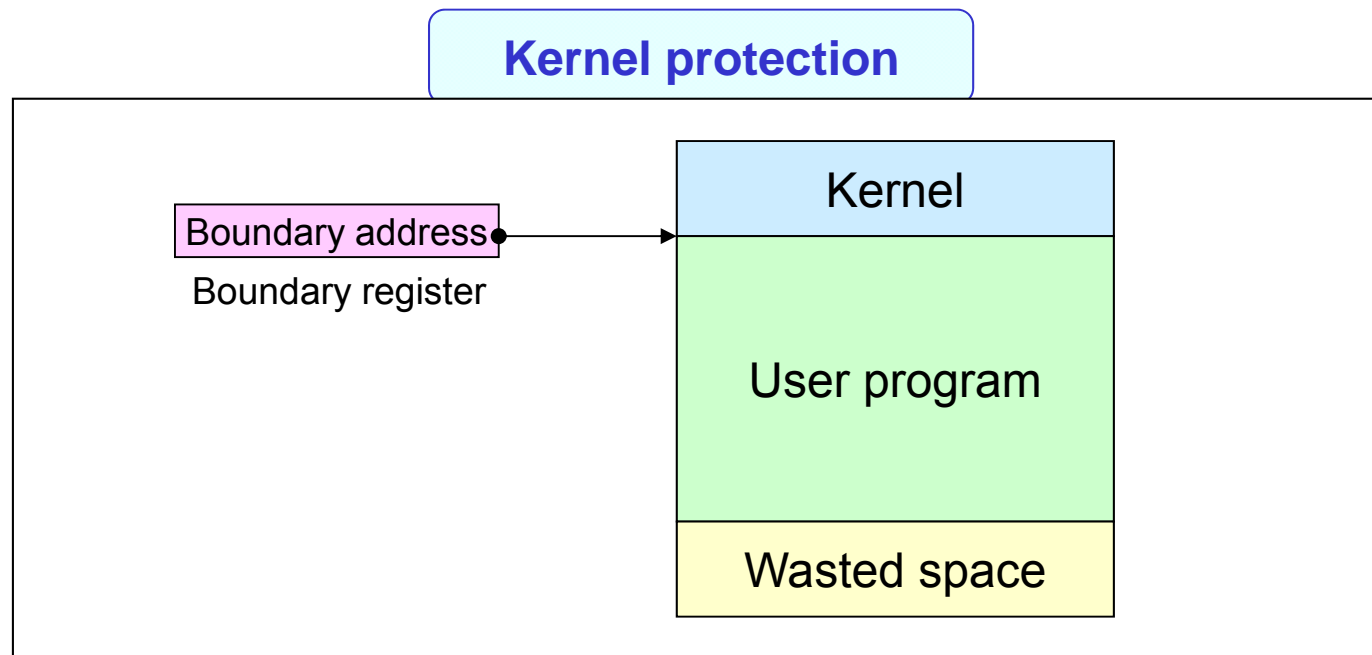


Uniprogramming

❑ Issue-2

❑ Kernel protection

✓ Boundary register





Uniprogramming

❑ Issue-3

- ❑ Low system resource utilization

- ❑ Low system performance

- ❑ Solution

 - ✓ Multiprogramming



FPM

❑ FPM: Fixed Partition Multiprogramming

- ❑ Divide memory into several fixed-size partitions
- ❑ One process in one partition
- ❑ When number of partitions = k
 - ✓ Maximum multiprogramming degree = k
- ❑ IBM OS/360 MFT



FPM

□ FPM Example

FPM Example

0	Kernel
a1	partition-A (10MB)
a2	partition-B (10MB)
a3	partition-C (20MB)
a4	partition-D (30MB)
a5	partition-E (50MB)



FPM

□ Data structure for FPM: Example

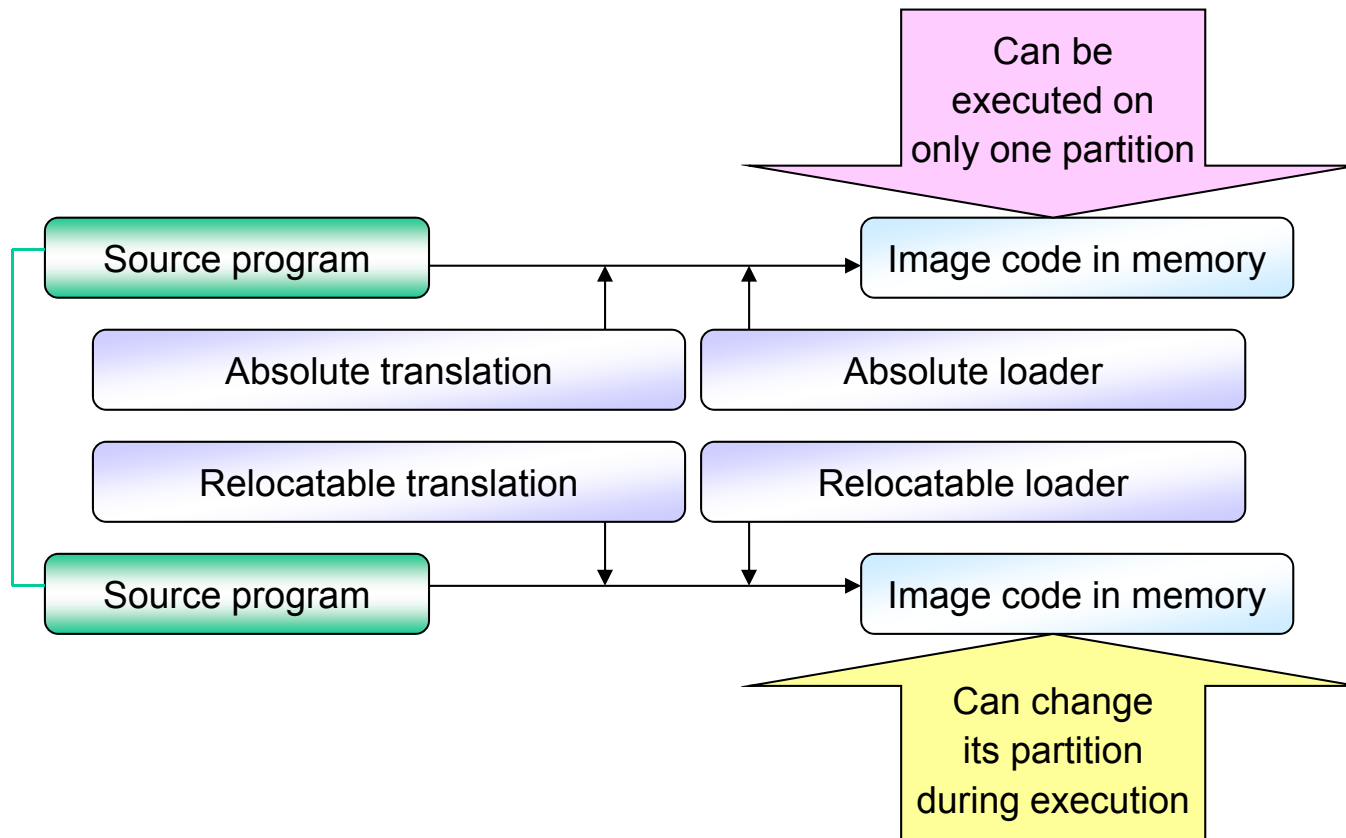
Data structure for FPM

partition	start address	size	current process ID	other fields
A	a1	10 MB	-	...
B	a2	10 MB	-	...
C	a3	20 MB	-	...
D	a4	30 MB	-	...
E	a5	50 MB	-	...



FPM

□ Program relocation

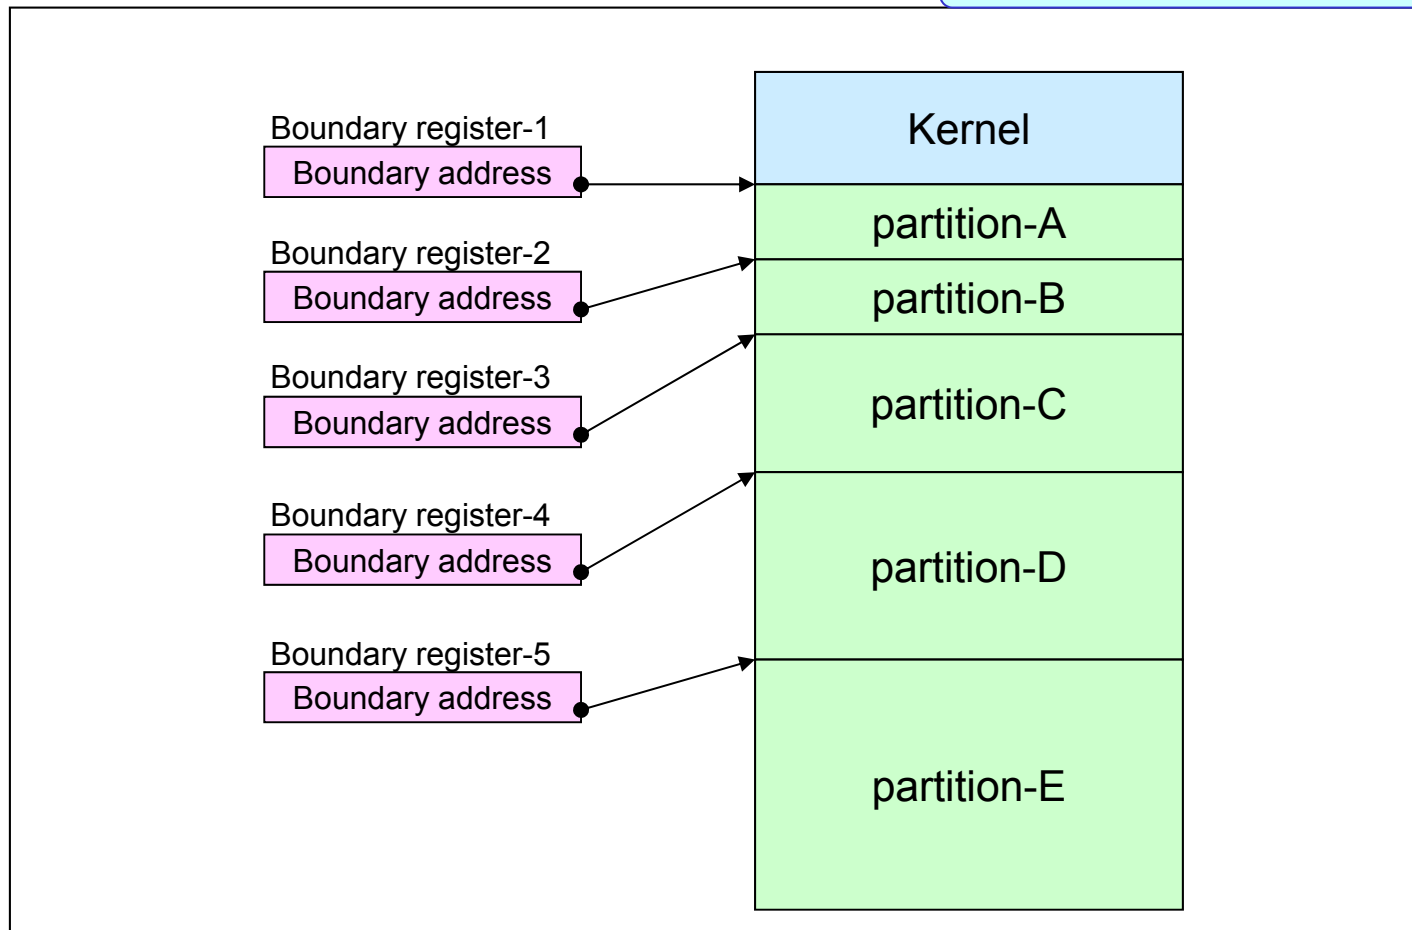




FPM

□ Protection method 1

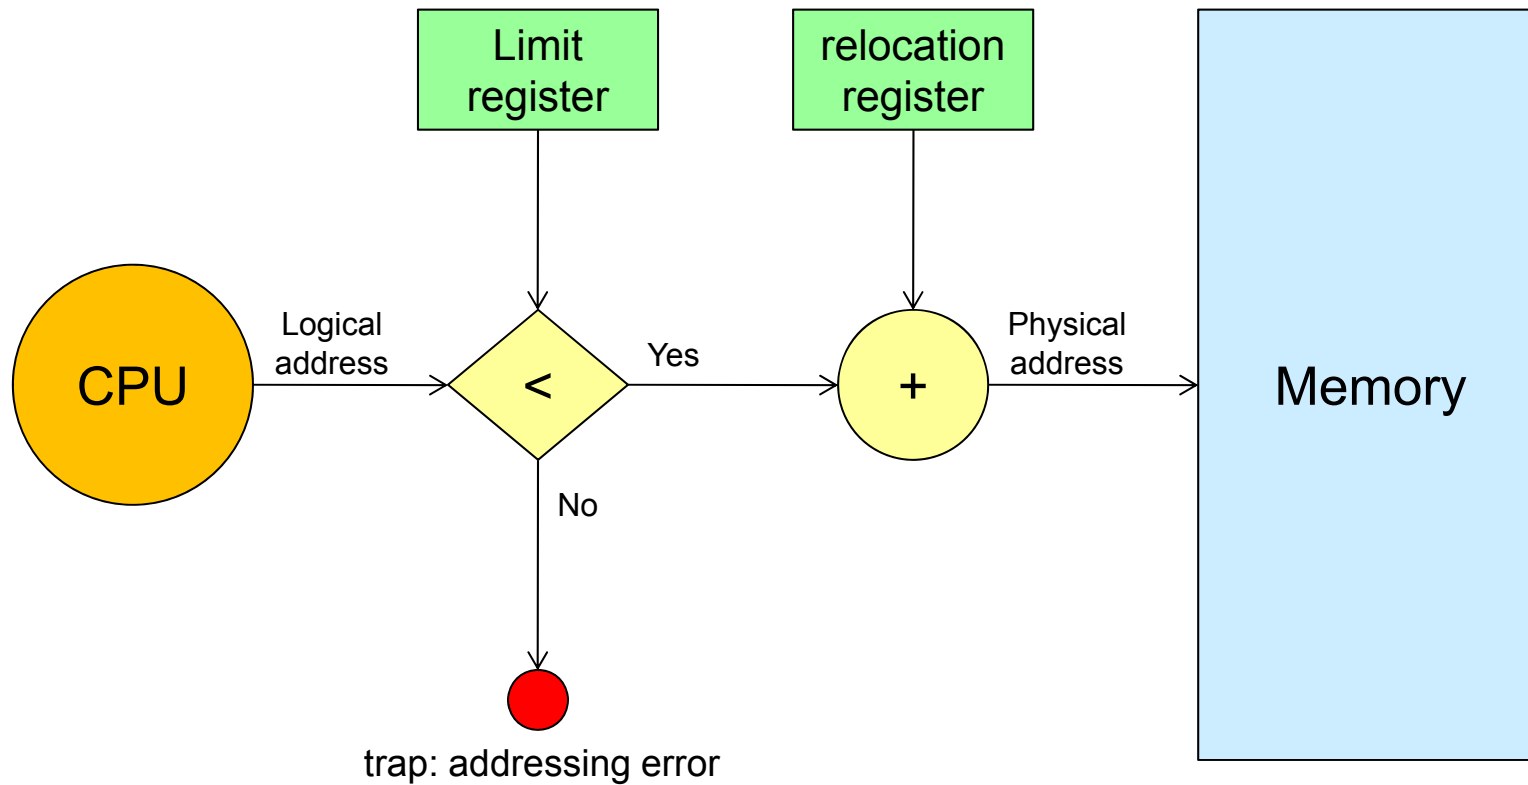
Protection in FPM





FPM

Protection method 2





FPM

□ Protection method 2

□ Loading the relocation and limit registers

- ✓ During context switching
- ✓ By the dispatcher
- ✓ Using a special privileged instruction
 - Allows the OS to change the value of the registers
 - Prevents user programs from changing the registers



FPM

❑ Fragmentation

❑ Storage space waste

✓ Internal fragmentation

- Exists when the memory space allocated is larger than the requested memory

✓ External fragmentation

- Exists when enough total memory space exists to satisfy the request, but it is not contiguous

❑ Both types of fragmentation exist in FPM



FPM

□ Summary

- Several fixed partitions in memory
- Simple memory management
- Low memory management overhead
- May cause poor resource utilization
- Internal/external fragmentation



VPM

❑ VPM: Variable Partition Multiprogramming

- ❑ Initially, all memory is available as a single large block of available memory (hole, partition)
- ❑ Memory partition state dynamically changes as a process enters (or exits) the system
- ❑ No internal fragmentation in a partition
- ❑ Contiguous allocation



VPM

□ Example

Memory allocation and partition scenario


♦ Assumption

- Memory space: 120 MB
- (a) : Initial state
- (b) : After loading process A(20MB)
- (c) : After loading process B(10MB)
- (d) : After loading process C(25MB)
- (e) : After loading process D(20MB)
- (f) : After process B releases memory
- (g) : After loading process E(15MB)
- (h) : After process D releases memory



VPM

VPM Example




partition	start address	size	PID	other field
1	u	120	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	100	none	...



VPM

VPM Example




partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	100	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	90	none	...



VPM

VPM Example




partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	90	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	25	C	...
4	u+55	65	none	...



VPM

VPM Example




partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	25	C	...
4	u+55	65	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	45	none	...



VPM

VPM Example




partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	B	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	45	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	none	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	45	none	...



VPM

VPM Example




partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	none	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	45	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	none	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	15	E	...
6	u+90	30	none	...



VPM

VPM Example



partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	none	...
3	u+30	25	C	...
4	u+55	20	D	...
5	u+75	15	E	...
6	u+90	30	none	...

partition	start address	size	PID	other field
1	u	20	A	...
2	u+20	10	none	...
3	u+30	25	C	...
4	u+55	20	none	...
5	u+75	15	E	...
6	u+90	30	none	...



VPM

❑ Placement strategies

❑ First-fit

- ✓ Start searching at the beginning of the state table
- ✓ Allocate the first partition that is big enough
- ✓ Simple and low overhead

❑ Best-fit

- ✓ Search the entire state table
- ✓ Allocate the smallest partition that is big enough
- ✓ Long search time
- ✓ Can reserve large size partitions
- ✓ May produce many small size partitions
- ✓ External fragmentation



VPM

❑ Placement strategies

❑ Worst-fit

- ✓ Search the entire state table
- ✓ Allocate the largest partition available
- ✓ May reduce the number of small size partitions

❑ Next-fit

- ✓ Similar to first-fit method
- ✓ Start searching from where the previous search ended
- ✓ Circular search the state table
- ✓ Uniform use for the memory partitions
- ✓ Low overhead



VPM

❑ Coalescing holes

- ❑ Merge adjacent free partitions into one large partition

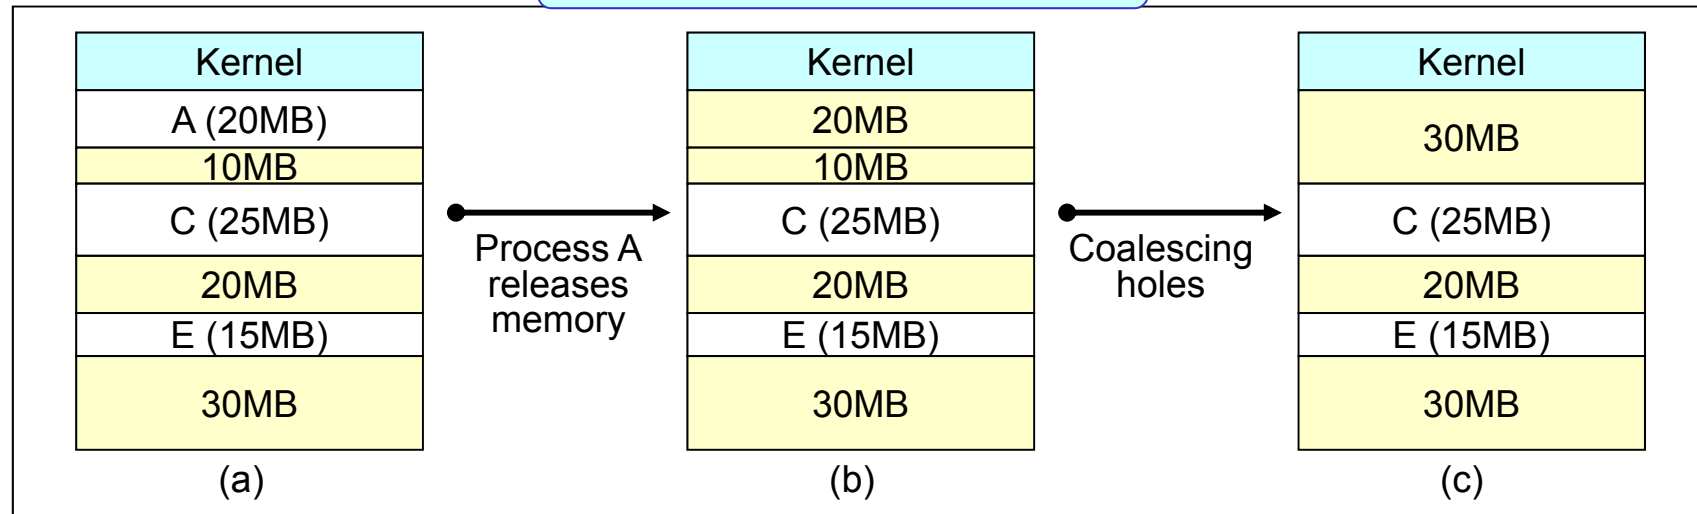
❑ Storage compaction

- ❑ Shuffle the memory contents to place all free memory together in one large block (partition)
- ❑ Done at execution time
 - ✓ Can be done only if relocation is dynamically possible
- ❑ Consumes so much system resources
 - ✓ Consumes long CPU time



VPM

Coalescing holes (1)



partition	start address	size	current process ID
1	u	20	A
2	u+20	10	none
3	u+30	25	C
4	u+55	20	none
5	u+75	15	E
6	u+90	30	none

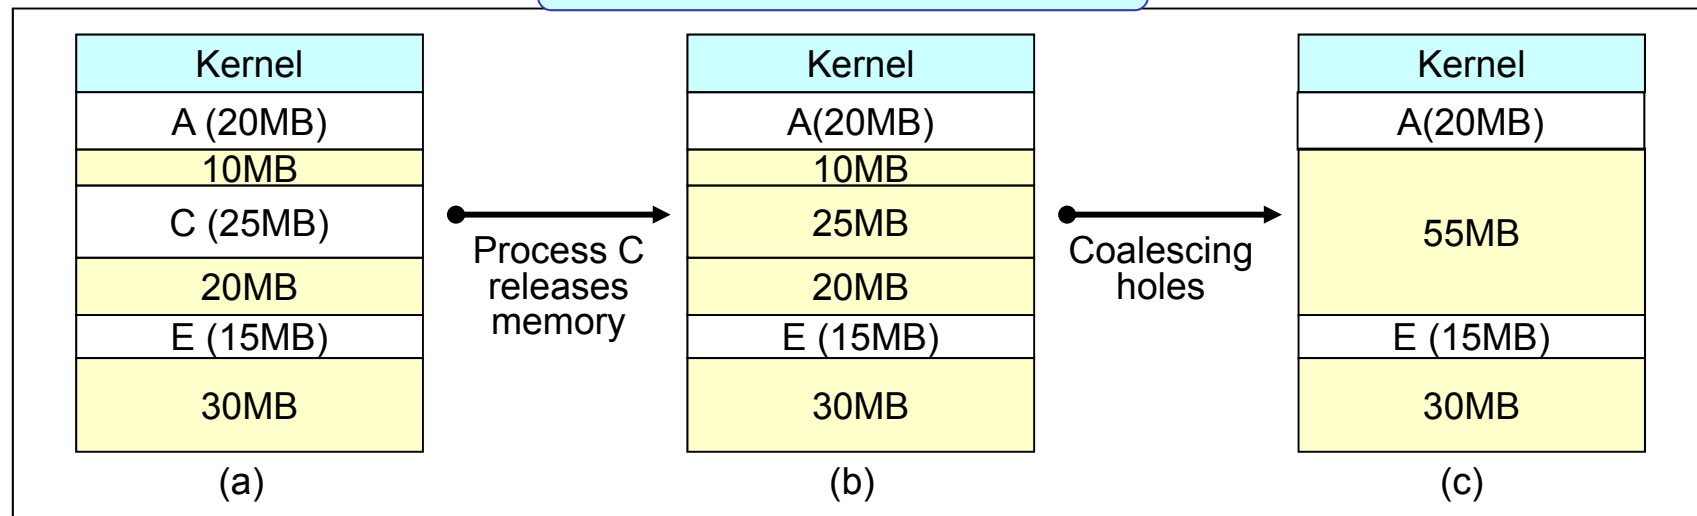
partition	start address	size	current process ID
1	u	20	none
2	u+20	10	none
3	u+30	25	C
4	u+55	20	none
5	u+75	15	E
6	u+90	30	none

partition	start address	size	current process ID
1	u	30	none
2	u+30	25	C
3	u+55	20	none
4	u+75	15	E
5	u+90	30	none



VPM

Coalescing holes (2)



partition	start address	size	current process ID
1	u	20	A
2	u+20	10	none
3	u+30	25	C
4	u+55	20	none
5	u+75	15	E
6	u+90	30	none

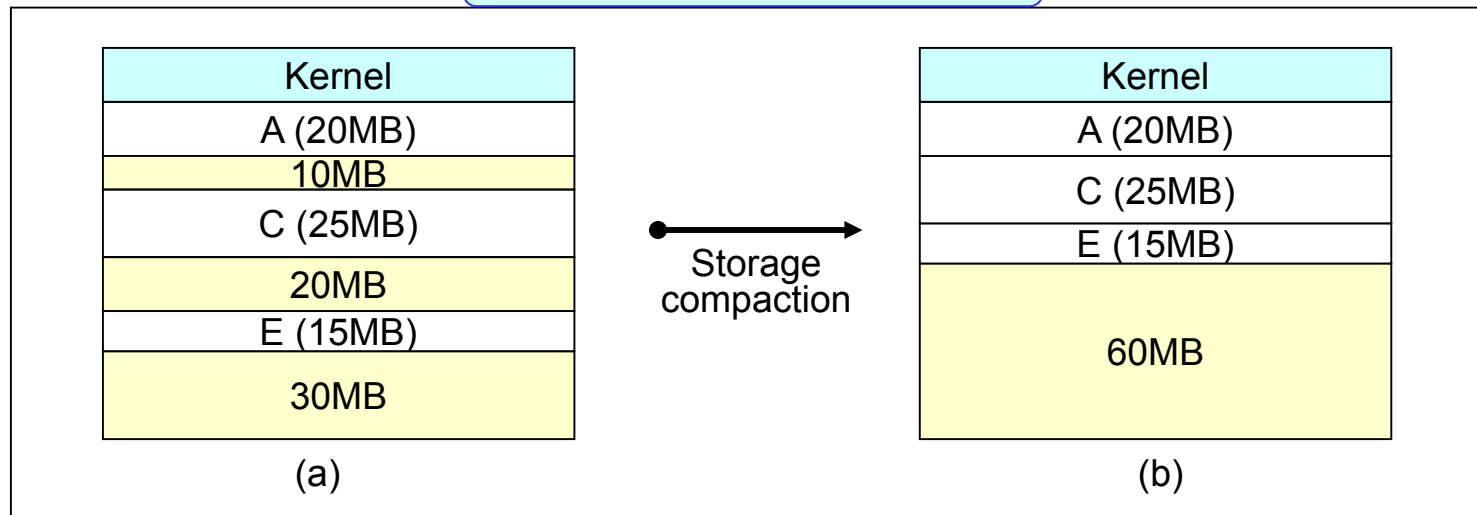
partition	start address	size	current process ID
1	u	20	A
2	u+20	10	none
3	u+30	25	none
4	u+55	20	none
5	u+75	15	E
6	u+90	30	none

partition	start address	size	current process ID
1	u	20	A
2	u+20	55	none
3	u+75	15	E
4	u+90	30	none



VPM

Storage compaction



partition	start address	size	current process ID
1	u	20	A
2	u+20	10	none
3	u+30	25	C
4	u+55	20	none
5	u+75	15	E
6	u+90	30	none

partition	start address	size	current process ID
1	u	20	A
2	u+20	25	C
3	u+45	15	E
4	u+60	60	none



Discontiguous Memory Allocation



Discontiguous Memory Allocation

- ❑ Paging
- ❑ Segmentation



Paging

□ Paging


- Memory management scheme that permits the physical address space of a process to be noncontiguous
- Implemented by closely integrating the hardware and operating system



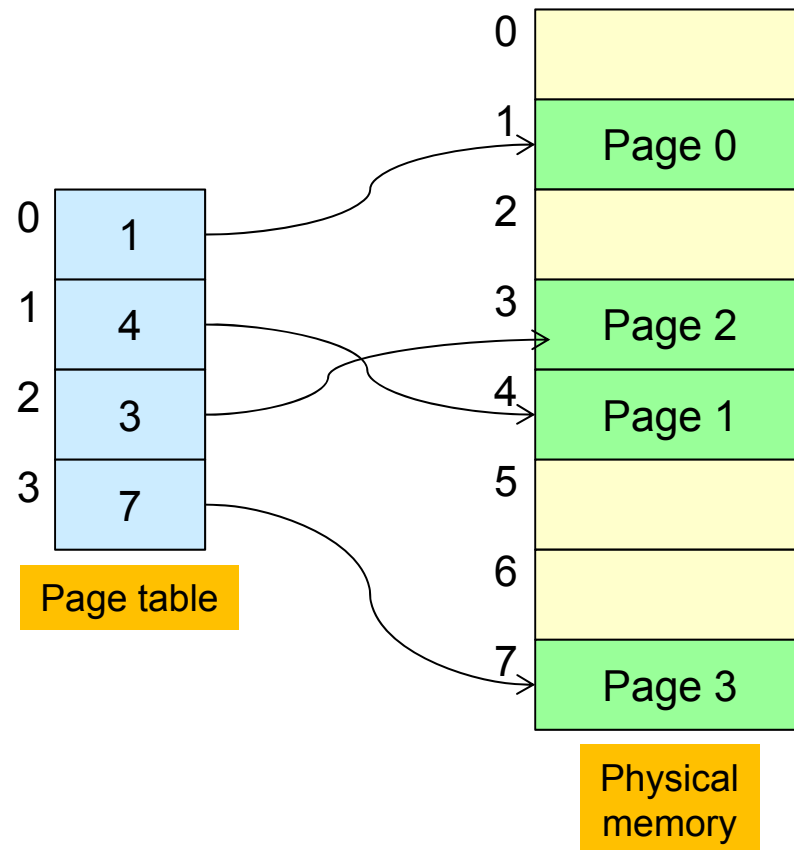
Paging

□ Paging

□ Example



Logical
memory





Paging

□ Basic method

□ Frames (page frames)

- ✓ Breaking physical memory into fixed-size blocks

□ Pages

- ✓ Breaking logical memory into blocks of the same size

✓ Page size

- Typically power of 2
- 512B ~ 16MB (typically 4KB ~ 8KB)

□ Logical address

✓ $\mathbf{v} = (\mathbf{p}, \mathbf{d})$

- **p**: page number
- **d**: page offset (displacement)



Paging

□ Basic method

□ Address mapping

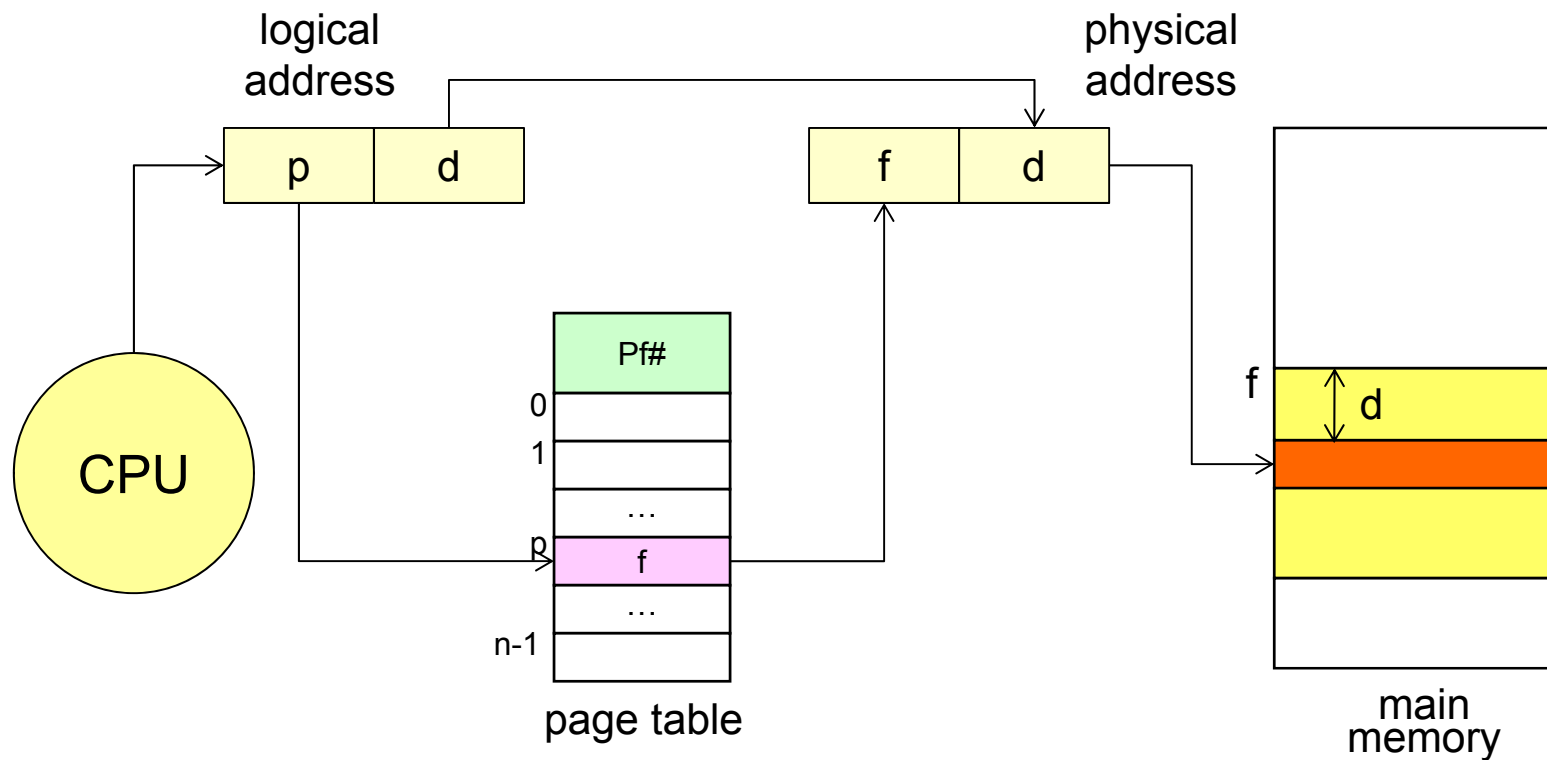
- ✓ Logical address → physical address
- ✓ Hidden from the user and controlled by the operating system
- ✓ Uses page table
 - A page table for each process



Paging

□ Basic method

□ Paging hardware and address mapping





Paging

□ Basic method

□ Page table implementation

- ✓ Dedicated CPU registers
- ✓ TLB (Translation Look-aside Buffer)

See [Virtual Memory Organization]



Segmentation

❑ Basic concept

- ❑ View memory as a collection of variable-sized segments
- ❑ View program as a collection of various objects
 - ✓ Functions, methods, procedures, objects, arrays, stacks, variables, and so on

❑ Logical address

- ✓ $\mathbf{v} = (\mathbf{s}, \mathbf{d})$
 - \mathbf{s} : segment number
 - \mathbf{d} : offset (displacement)

segment-0
segment-1
segment-2
segment-3
segment-4



Segmentation

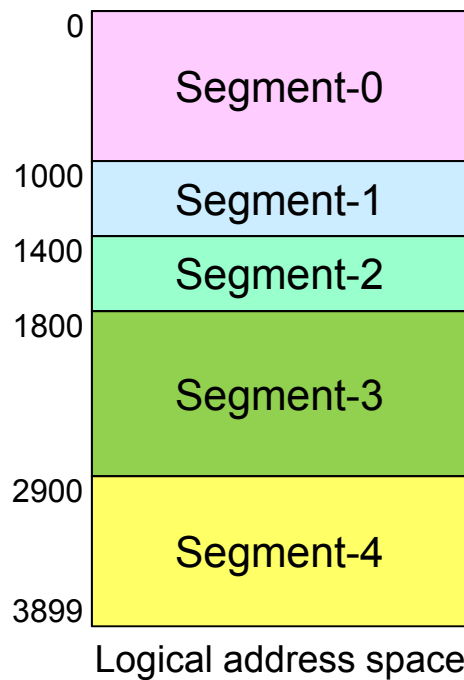
□ Basic concept

- Normally, when the user program is compiled, the compiler automatically constructs segments reflecting the input program
 - ✓ Code
 - ✓ Global variables
 - ✓ Heap
 - ✓ Stacks
 - ✓ Standard library



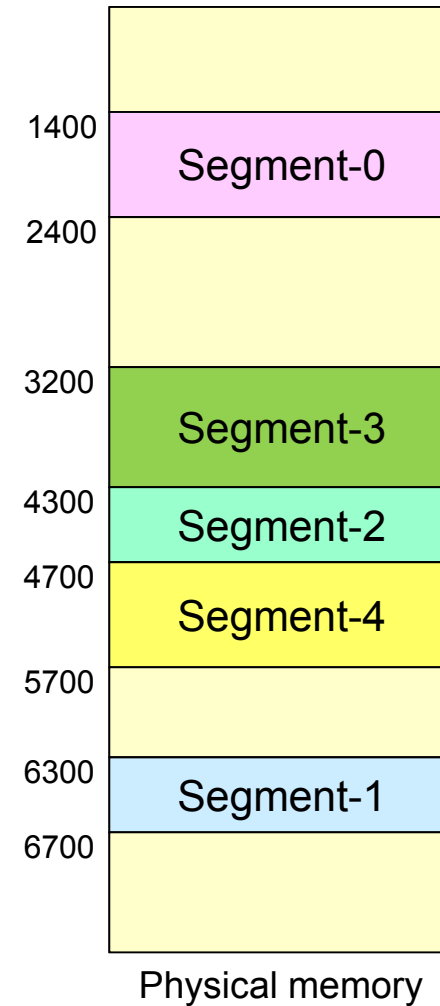
Segmentation

□ Example



	base	limit
0	1400	1000
1	6300	400
2	4300	400
3	3200	1100
4	4700	1000

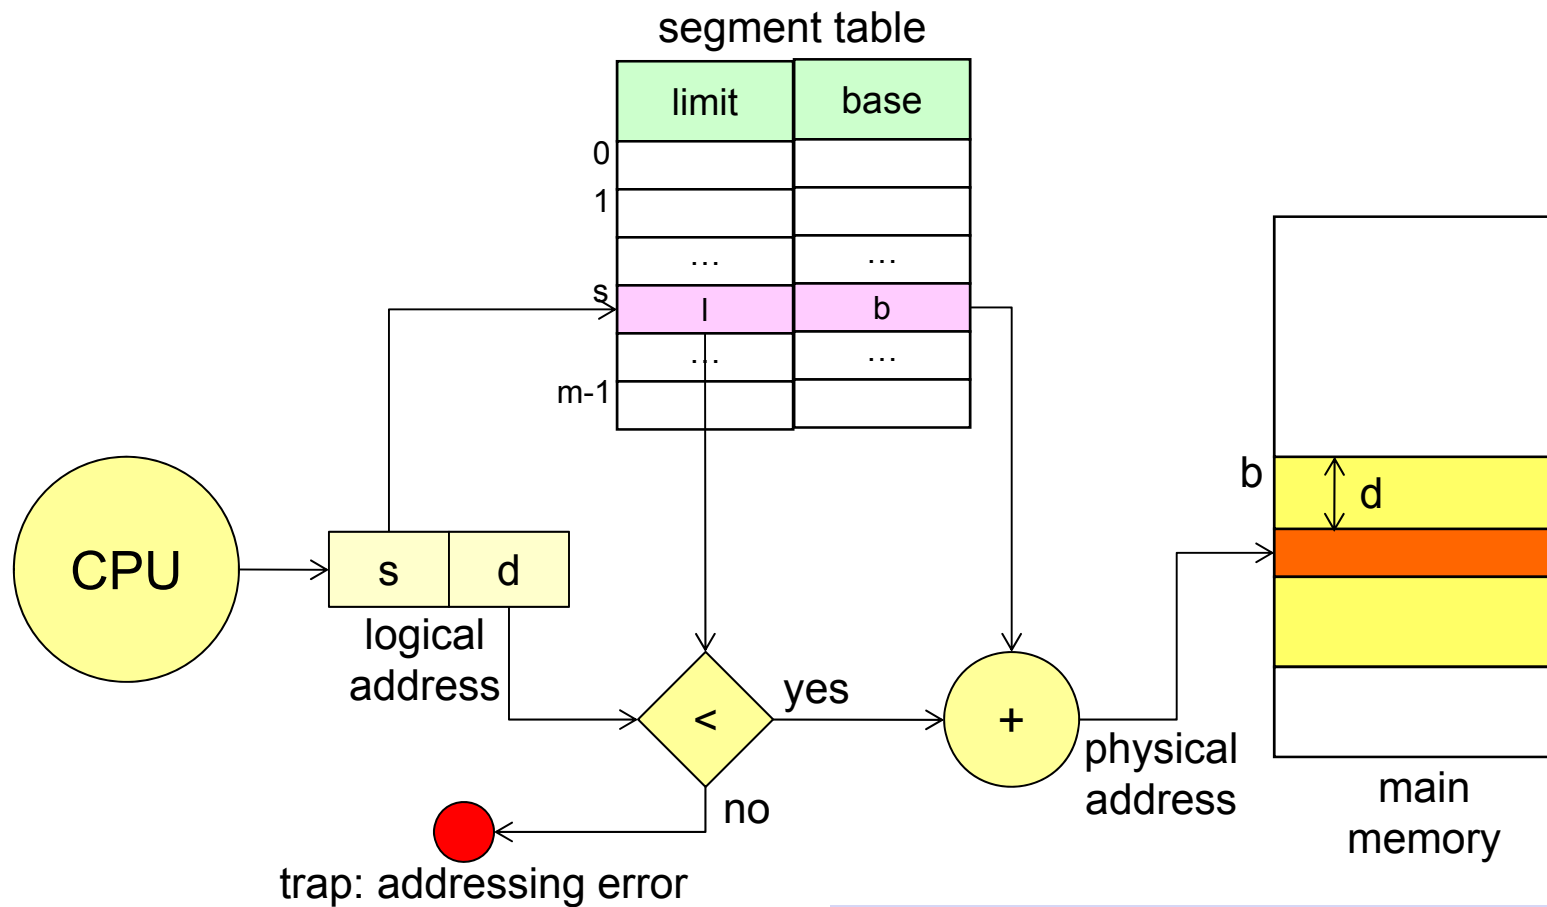
Segment table





Segmentation

□ Address mapping





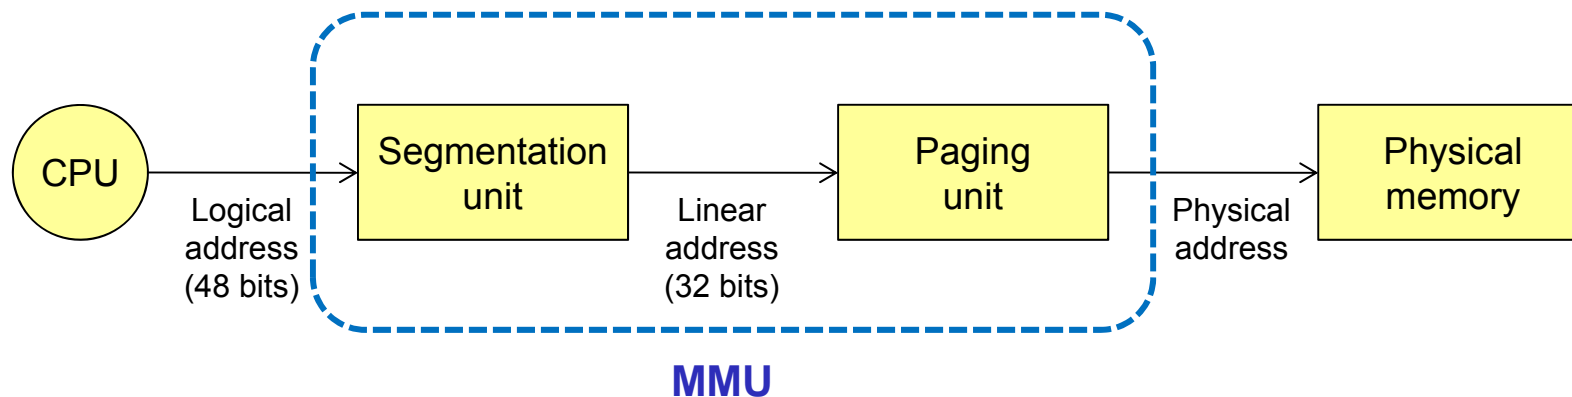
Intel Pentium Architecture



Intel Pentium Architecture

Intel Pentium architecture

- Supports both pure segmentation and segmentation with paging
- Address mapping





Intel Pentium Architecture

❑ Pentium segmentation

- ❑ Segment size: max 4GB
- ❑ Max number of segments per process: 2^{14}
- ❑ Logical address space of a process
 - ✓ 1st partition: 2^{13} private segments
 - LDT(Local Descriptor Table) keeps the information
 - ✓ 2nd partition: 2^{13} shared segments
 - GDT(Global Descriptor Table) keeps the information
- ❑ Each entry of LDT and GDT
 - ✓ 8-byte segment descriptor
 - ✓ Detailed information on a particular segment
 - Base location, limit of the segment, etc



Intel Pentium Architecture

□ Pentium segmentation

□ 6 segment registers

- ✓ Allows 6 segments to be addressed at any time
- ✓ Points to the appropriate entry in LDT or GDT

□ 6 8-byte micro-program registers

- ✓ Corresponding descriptors from either the LDT or GDT




Intel Pentium Architecture

❑ Pentium segmentation

❑ Logical address: (**selector**, **offset**)

✓ **selector**: (**s**, **g**, **p**)

- **s**: segment number (13 bits)
- **g**: whether the segment is in the LDT or GDT (1 bit)
- **p**: deals with protection (2 bits)



✓ **offset**

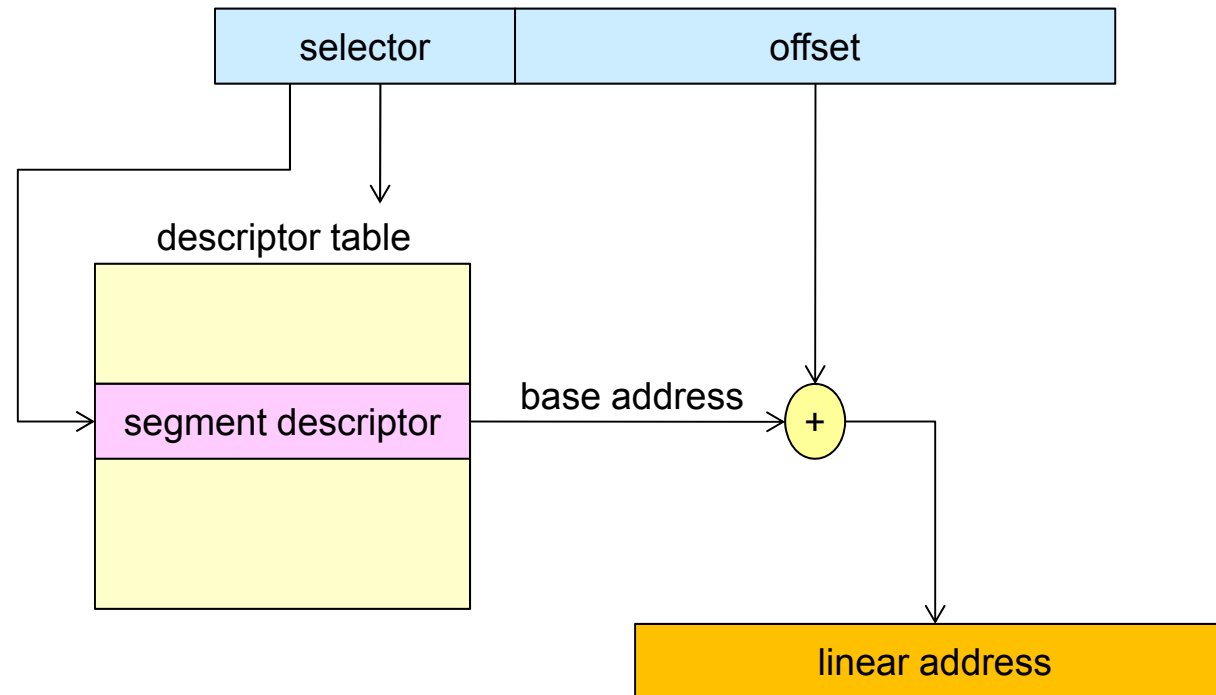
- 32-bit number (location within the segment)



Intel Pentium Architecture

□ Pentium segmentation

□ Linear address generation






Intel Pentium Architecture

❑ Pentium paging

- ❑ Page size: 4KB or 4MB


- ❑ For 4KB pages

 - ✓ Two-level paging scheme is used



- ❑ For 4MB pages

 - ✓ Single-level paging scheme is used

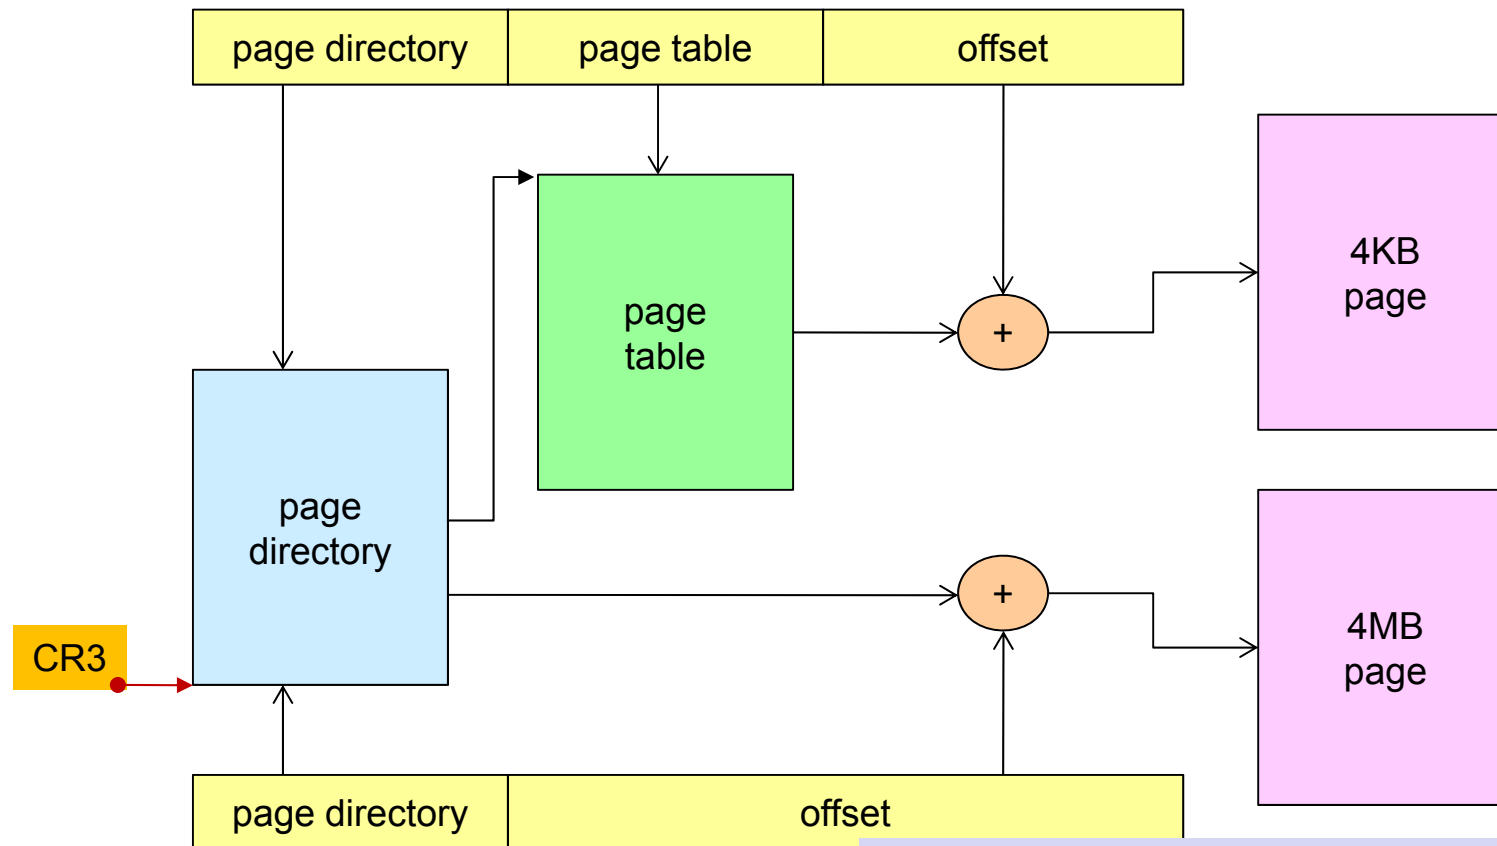




Intel Pentium Architecture

❑ Pentium paging

❑ Physical address generation



A. Silberschatz, et. al., Operating System Concepts, 8-ed., Wiley, 2010.



Intel Pentium Architecture

□ Pentium paging

□ Notes

- ✓ **Page size** flag in the page directory entry
 - Indicates whether the size of the page frame is 4KB or 4MB
- ✓ **Invalid** bit in the page directory entry
 - Indicates whether the page table to which the entry is pointing is in memory or on disk
 - When the table is on disk, the OS can use the other 31 bits to specify the disk location of the table
 - The table is brought into memory on demand



Intel Pentium Architecture

❑ Linux on Pentium systems

- ❑ Linux does not rely on segmentation and uses it minimally
- ❑ Linux uses only 6 segments
 - ✓ Kernel code segment
 - ✓ Kernel data segment
 - ✓ User code segment
 - ✓ User data segment
 - ✓ Task state segment (TSS)
 - ✓ Default LDT segment



Intel Pentium Architecture

❑ Linux on Pentium systems

- ❑ User code and user data segment
 - ✓ Shared by all processes running in user mode
- ❑ One TSS for each process
 - ✓ Used to store H/W register context during context switches
- ❑ Default LDT segment
 - ✓ Generally not used



Intel Pentium Architecture

❑ Linux on Pentium systems

❑ 2-bit protection field in segment selector

- ✓ Allows 4 levels of protection
- ✓ But, Linux uses only 2 modes
 - User mode and kernel mode

❑ Linux adopted 3-level paging scheme

- ✓ Works well in both 32-bit and 64-bit architectures
- ✓ Linear address

global directory	middle directory	page table	offset
---------------------	---------------------	---------------	--------

- ✓ For Pentium systems, the size of the middle directory is zero bits



Summary

☐ Contiguous memory allocation

- ☐ Uniprogramming
- ☐ Multiprogramming
 - ✓ FPM
 - ✓ VPM

☐ Discontiguous memory allocation

- ☐ Paging
- ☐ Segmentation

☐ Intel Pentium system

- ☐ Paging and segmentation
- ☐ Linux on Pentium systems